CoSMo: Intent-based Composition of Shader Modules

Georg Haaser, Harald Steinlechner, Michael May
Michael Schwärzler, Stefan Maierhofer and Robert Tobler
haaser@vrvis.at

VRVis Research Center, Donau-City-Strasse 1, Vienna, Austria

Keywords: Shader, Composition, Rendering, Language, Embedded

Abstract: We propose a novel shader programming model which operates on *intent-oriented shader modules* instead of specific programs for dedicated GPU rasterization pipeline stages. In constrast to existing pipeline shader frameworks, our system exposes a radically simplified pipeline, which we purposefully aligned with our basic intuition of shaders as per-primitive and per-pixel operations and compositions thereof. This simplicity lends itself to structure modules purely based on their intent, instead of dealing with structure enforced by specific versions of graphics APIs. Consequently, this offers great flexibility when it comes to reusing and combining modules with completely different semantics, or when targeting different graphics APIs. The simplicity and uniformity of our system also motivates automatic parameterization and simplification of shader programs as well as interesting interactive shader development and management techniques.

1 Introduction

Programming shaders for hardware-accelerated rasterization frameworks like DirectX or OpenGL has attained an important role in the development process of rendering applications. Even though the flexibility and possibilities in graphics development have drastically improved with the introduction of these shaders over the last few years, recent advances in "CPU-based" programming languages and software engineering are often not reflected in shader-programming. Especially the limitations in terms of shader management in larger software projects cause the tasks of *combining shader effects*, *targeting different hardware*, *supporting older API versions* or *optimizing these shader permutations* to become extremely time-consuming, tiresome, and error-prone.

The C-style definition of a single shader stage program is only simple during the primary creation process: As soon as such an effect has to be *combined* with other shaders to generate the desired final surface illumination, or has to be used on another API version or target platform, programmers either tend to build large, complex Über-Shader constructs with computationally expensive dynamic branching techniques, or manage hundreds of shader combinations and permutations manually. Object-oriented approaches (McCool et al., 2002; Kuck and Wesche, 2009; Foley and Hanrahan, 2011) and novel Shader Model 5 functionality (e.g. `interfaces` (Microsoft, 2010)) extend procedural languages with abstractions like interfaces and limit code duplication via inheritance. Inheritance as mechanism for composition however has shortcomings in terms of ad hoc compositions and reusability, as each composition has to be stated explicitly (see sections 2.1 and 2.2) and in terms of extensibility, as extension points have to be anticipated by providing abstract or virtual methods.

Based on these observations, we propose a novel shader programming model that emphasizes the *intent* of a shader, based on the following ideas:

- *abstract shader stages*: by freeing the shader modules from concrete pipeline stages, we let the programmer specify what he wants his shader modules to do in an abstract, backend-independent manner
- *composition via semantic input/output types*: with the introduction of semantically annotated input and output types, these types encode the intent of what is computed by each shader module, and thus composition operators can be built, that combine the modules according to this intent
- *fine-tuning of semantics*: by providing more detailed information for the semantic types such as computation rates, the programmer can exactly specify the intent of his shader

Using these ideas to provide a system that au-

Figure 1: Example built with composed shader modules (from left to right): transformation and per-pixel lighting, transformation and texturing, transformation/normal mapping/texturing and lighting, transformation/normal mapping/texturing/point sprite generation and lighting, transformation/normal mapping/texturing/point sprites/thick line generation and lighting.

tomatically combines modules based on their intent, we overcome the combinatorial explosion of typical shader systems where each and every combination has to be specified explicitly: Modules are typically expressed only in their most general form, and can be composed either statically as hard-coded expressions, or programmatically, which is useful to generate shaders based on runtime information or whole families of related shaders. Our high-level shader-code requires a specific functionality to be defined only once—no matter how often it is combined with other shaders and on how many target platforms it is deployed—while unneeded calculations are automatically eliminated. We leave the error-prone task of finding the optimal shader stage for each computation to the machine, which automatically maps shaders onto specific pipeline architectures (e.g. DirectX), performs global and local optimizations and code generation for distinct shader permutations, and finally emits a low-level shader program (e.g. HLSL) comparable to hand-crafted code.

2 Background

The ancestors of today's shading languages are Cook's shade trees (Cook, 1984) and Perlin's image synthesizer (Perlin, 1985). Cook's shade trees classify independent aspects like lighting, surface and volume into separate modules called shading processes. As a mechanism for composition each process is represented as an expression tree which supports grafting of commonly used expressions into other processes. However, the underlying model of computation which is purely declarative allows for no conditional control flow like loops as well as no mutable state. Perlin's image synthesizer is based on imperative procedures and therefore dissolves these limitations, but abandons the idea of logically independent shading processes. Procedures work on streams of fragments, and describe shading computations after hidden surface removal.

The most prevalent shader languages for real-time rendering (Cg (Mark et al., 2003), HLSL (Microsoft, 2012), and GLSL (Kessenich et al., 2012)) follow the *shader-per-stage* approach. Similarly to Perlin's image synthesizer each stage works on streams of objects like *vertices*, *primitives* or *fragments*. As a consequence they directly reflect the various pipeline stages of the hardware in the language itself. Although there are little restrictions in terms of algorithms that can be formulated, a corresponding shader function must be provided for each of the stages.

2.1 Towards composability

Shader languages like HLSL provide procedures as their main structuring mechanism. Über-Shaders usually implement the sum of all desired features and use ad-hoc mechanisms like macros and plain text processing for specialization and feature selection.

Metaprogramming frameworks (McCool et al., 2002; McCool and Toit, 2004; McGuire, 2005; Kuck and Wesche, 2009) overcome the lack of language level abstractions by utilizing meta-programming and macros. LibSh (McCool et al., 2002) provides an embedded language in C++, utilizing its features like objects and templates for combining shaders. McCool et al. (McCool et al., 2004) extends LibSh with algebraic combinators *connection* and *combination* which provides an expressive basis for combining shader functions.

Elliot (Elliott, 2004) proposes Vertigo, an embedded domain specific language written in Haskell that provides combinators in a very natural way. Based on these combinators, an implementation of a sophisticated shading infrastructure comparable to RenderMan Shading Language (RSL) (Hanrahan and Lawson, 1990) is demonstrated, including a subsequent compilation process which creates vertex- and fragment-shader programs.

Abstract shade trees (McGuire et al., 2006) are

based on a visual programming approach for shaders, and also provide automatic linkage of shader parameters as well as semantic operations like vector basis conversion. Although different shader components compose well, geometry shaders and tessellation are not treated at all. Trapp et al. (Trapp and Döllner, 2007) structures GLSL shader code into code fragments, each typed with predefined semantics. Code fragments may be composed at run-time and compiled to Über-Shaders. Of course Über-Shaders suffer from bad performance. Like other metaprogramming approaches the system cannot provide proper semantic analysis and cross-fragment optimization.

2.2 Towards pipeline shaders

The RenderMan Shading Language by Hanrahan and Lawson (Hanrahan and Lawson, 1990) combines the expressiveness of Perlin's image synthesizer with independent shader processes introduced by Cook. The concept resembles object-oriented classes, whereby each virtual method corresponds to an entry point called by the render system. Subclasses like *surface*, *light* and *volume* may be attached to surfaces. Furthermore RSL extends the concept with computation rates, i.e. the notion of inputs varying two different rates: `uniform` and `varying`. Specialized control-flow constructs provide mechanisms for communication between shaders.

A further refinement for computation rates was introduced by Proudfoot et al. (Proudfoot et al., 2001) in their Stanford Real-Time Shading Language (RTSL): `constant`, `primitive group`, `vertex` and `fragment`, where the last two rates directly corresponded to the stages of early programmable GPUs. Like Cook's shade trees (Cook, 1984), RTSL programs are purely declarative and can therefore be represented as DAGs, which affects expressiveness (e.g. limited data dependent control flow).

Renaissance (Austin, 2005) takes a more general approach and represents different shader pipeline stages as single functional shader programs. Parameters implicitly correspond to different computations rates. Compilation automatically lifts expressions into the earliest possible pipeline stage while maintaining semantics. However, Renaissance lacks support for structuring monolithic shader programs into well defined reusable modules, and no semantics for lifting expressions to groupwise shader stages (e.g. geometry shaders) are presented.

Foley and Hanrahan introduce Spark (Foley and Hanrahan, 2011), a *pipeline shader* approach based on RTSL (Proudfoot et al., 2001). Its two-layer approach uses declarative shader graphs on top of procedural subroutines and therefore combines the approaches of Cook (Cook, 1984) and Perlin (Perlin, 1985). Spark expands RSL's idea of treating a shader in an object-oriented way by using extending, virtual-, and abstract identifiers for compositing and customizing shaders. Rate-qualifiers and conversions between different rates are extensible and thus defined individually by each supported pipeline. Different modules may be composed by using mixin inheritance. Like other Über-Shader approaches before, Spark does not solve the combinatorial explosion problem because each composition must be stated explicitly.

3 Design

3.1 A shader as a pixel-valued function

Shader programming targets a highly parallel execution environment, where shading can be performed independently for each surface point, therefore functional programming is a natural match for specifying shaders (Cook, 1984). Although parts of a shader can be programmed in procedural style using local variables and loops, a complete shader program only has a single output value—the target pixel—and can thus be viewed as a single function. By using tail recursion instead of loops, and higher-order functions for control-flow it is even possible to map any procedural shader program to a purely functional representation. Rennaissance (Austin, 2005) is an example of such a functional approach to shader programming.

Since the output value of a shader expression for a single pixel can be an aggregate of multiple simple values (e.g. it can contain a colour, a depth value, etc.), we use the term *shader module* to denote a shader function with multiple input and output values. Multiple output values are programmatically handled by returning a single structure containing the individual output values.

Although our approach is based on the composition of such shader modules, and thus retains the expressiveness and extensibility of a functional design (which goes beyond what is possible with the specialised control flow elements introduced in RSL (Hanrahan and Lawson, 1990)), we have included control flow functions that are modelled on imperative languages in order to cater to shader programmers that are used to imperative shader languages. Details on these control flow functions are given in Section 4.

3.2 Semantic composition of shader modules

Combining shaders modules that are formulated as expressions can be done in a pipeline approach, by routing the output of one component into the input of another component. In order to derive the necessary composition functions for combining shader modules we will look at a simple example that combines three shader modules with two composition operators, namely *Sequence* and *Combine(operator)* that are applied to the individual output fields (see Figure 2). Note, that the different composition operators need to be applied to different types of input and output (*Normal*, *Color*, *LightPos*). We use the term *semantic* for these types, as they go beyond the typical notion of data types in a language: both *Normal* and *LightPos* are represented as float vectors, but this does not capture their intent.

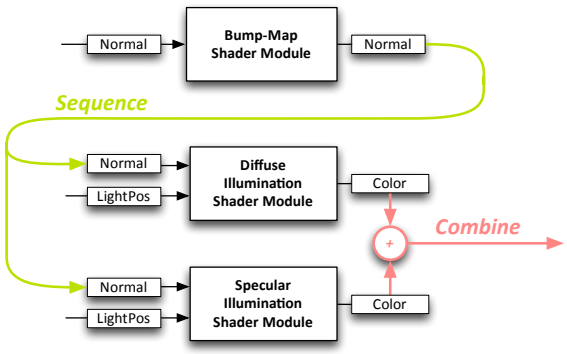

Figure 2: The output values of three simple shaders can be combined with two composition operators: *Sequence* and *Combine(operator +)*.

Generalizing from this example, we define the following three basic composition operators for semantic shader composition, that operate on arbitrary shader modules each with one or more semantic input types and one or more semantic output types (see Figure 3):

$Compose(module_1, module_2)$ composes the output of the two shader modules. All output semantics of the two input shader modules must be different.

$Sequence(module_1, module_2, \{semantic_1, ..., semantic_M\})$ combines the specified semantics of the supplied shader modules in sequence. The remaining output semantics of the two input shaders must be different.

$Combine(module_1, module_2, semantic, operator)$ applies the supplied binary operator to the output of the two input shaders with the given semantic to return a value of the same semantic. The remaining output semantics of the two input shader modules must be different.

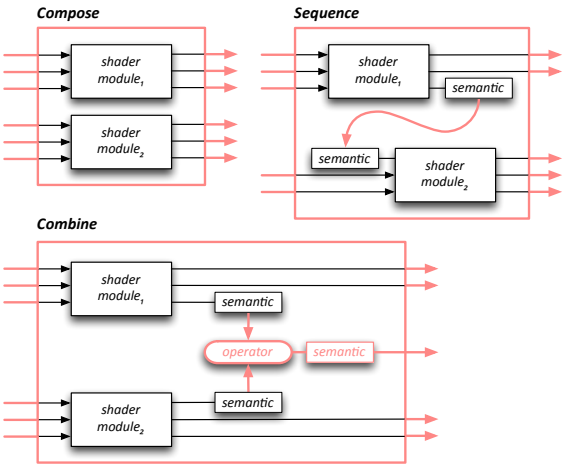

Figure 3: The three semantic-aware basic composition operators for shader composition. For illustration purposes, the *Sequence* function is depicted with only one semantic routed between its shader module arguments.

For all composition operators, the input semantics of the two input shader modules are allowed to be either partially or completely equal. In this case, the same value is supplied to both shaders. These basic composition operators add the concept of semantic-specific operations to the usual composition functions used in functional languages. On top of these basic composition operators we can now define a more general composition function that sequences, combines, and composes multiple shader modules based on their semantic:

$$\begin{aligned} Composition(module_1, \quad module_2, \quad \ldots \quad module_N, \\ semantic_1 \quad : \quad composition_1, \\ \vdots \qquad\qquad \vdots \\ semantic_M \quad : \quad composition_M, \\ default \quad : \quad composition_{default}) \end{aligned}$$

where a separate composition operator (either *Sequence* or $Combine(operator)$) is specified to combine each semantic and *Compose* is wrapped around the result.

We provide a number of convenience compositions in our approach, that are specializations of this general composition function with various predefined function and operator arguments. For convenience we also predefine simple shaders for changing semantics (e.g. $Position \rightarrow Color$). An example composition can be found in Section 5.

Our approach of automatically combining shaders based on their semantically tagged inputs and out-

puts is inherently more flexible than a static object-oriented approach as implemented by Spark (Foley and Hanrahan, 2011):

- The object-oriented way of extending functionality by overriding virtual functions requires, that each possible extension point needs to be foreseen by the implementer of the base shaders. Since only a limited number of possible ways of extending functionality can be provided in a typical design of such base shaders, the extensibility of such an object-oriented approach is necessarily limited.
- Due to the static way of combining and extending shaders, each and every new combination of simple shader functions must be explicitly and manually implemented. Since the number of combinations of simple shaders is exponential in the number of shaders, this leads to a combinatorial explosion that cannot be handled by a static approach. The use of a composition function as shown above, makes it possible to automatically combine simple shaders based on the geometry that needs to be rendered: the rendering framework can analyse the properties of the geometry, and combine only the simple shaders that are actually needed for rendering the combination of properties encountered.

All composition possibilities offered by a static object-oriented approach can be easily built using a sub-set of the available functionality in our meta-function approach:

- Each virtual method corresponds to a semantic tag: different simple shaders can perform different operations on the input with the same semantic tag. Changing the implementation of one virtual method thus corresponds to replacing one of the simple shaders in a composition of multiple shaders.
- The effect inheritance in the object-oriented approach can be realised using the *combine* composition operator on two simple shaders that correspond to the base-class implementation and the overriding implementation. By using a function that ignores the result of the simple shader corresponding to the base-class the result of the combination corresponds to the result of the overriding simple shader.

Thus our approach provides a superset of the functionality provided by the object-oriented approach, and the additional functionality eliminates the large number of shader combinations that have to be manually specified.

3.3 Abstract stages

The various stages in the shader pipeline can be viewed as optimizations on the single pixel-value function, in order to reduce the number of evaluations of various expressions (for an example see Figure 4).

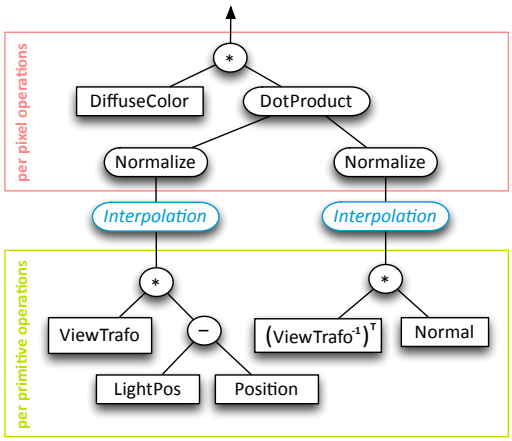

Figure 4: Optimizing the evaluation of a shader expression by evaluating the parts at different stages. The indicated interpolation functionality is provided by the hardware stages.

Although in principle, every shader could be formulated as a single function that returns a pixel value, this would require the implicit interpolation that is performed between the vertex and fragment stages of the shader pipeline (see Figure 4) to be explicitly specified in this function. In order to overcome this inconvenience, we propose to retain the notion of shader stages, but as opposed to the multiple hardware stages we only specify two abstract stages, that turn out to be sufficient in practice :

per-pixel operations :
also called per-fragment operations, these are all the operations that need to be performed for each pixel or fragment. Typically this includes all traditional shading operations that affect the material of an object. In functional notation, these operations perform the mapping: *per pixel parameters* $\rightarrow$ *per pixel output*.

per-primitive operations :
all operations performed for each primitive (e.g. triangle or line). This typically encompasses geometric transformations. In functional notation, these operations perform the following mapping: *per primitive parameters* $\rightarrow$ *per primitive output*.

Thus all the simple shaders that can be composed in our framework consist of explicitly specified per-pixel and/or per-primitive operations, and thus each of our simple shaders can be viewed as either a partial

or a fully specified, but still abstract pipeline of operations (see Figure 5). Since we do not explicitly specify operations for a specific hardware pipeline, all our shaders are still specified in an abstract manner, and need to be explicitly mapped onto the hardware stages of a concrete pipeline.

This has some consequences for the composition operators defined in the previous section: specifically, a per-pixel output of a shader module cannot serve as an input for a per-primitive shader module. If each shader module that consisted of a per-geometry stage and a per-pixel stage was viewed as a monolithic block with no allowed change in the routing of data between the per-geometry and the per-pixel stage, this would lead to significant limitations on which shader-modules could be combined.

In order to avoid that, we combine the stages of our shader modules individually, i.e. the composition works independently on the per-primitive stage and on the per-pixel stage. This makes it possible to combine shader modules as simply as if they were single stage modules, and retain the intended functionality programmed in the different stages. The semantically tagged inputs and output of the individual stages are available for automatic composition with other shader modules (see Figure 5).

Thus we extend our concept of shader modules to encompass whole shader pipelines, on which our composition operators work, retaining the semantic input and output types between the abstract stages, which is beyond the functionality of the algebraic combinators of McCool et al. (McCool et al., 2004).

3.4 Optimization

Naïvely mapping such high-level shaders onto the hardware stages of a concrete pipeline such as the DirectX 11 pipeline results in a lot of overhead due to a number of possible inefficiencies. In order to provide comparable performance in our system, we perform several optimisations highlighted in the following section.

Dead code elimination. Since our modules are programmed to be maximally reusable, they are implemented to cover the most general case, and thus provide a large number of semantic outputs which can be used by other modules that are placed later in a composition. Thus it is of vital importance for the overall performance to identify all unused outputs, and eliminate these from the composition result: this is done by starting with all used pixel outputs and tracking all necessary inputs back through the pipeline. All unused inputs and outputs are removed, essentially performing a dead code elimination step.

Backend stage mapping. A typical hardware shader pipeline has a number of stages that can be used to perform the operations in our pipelined shaders. Based on the abstract stage, the following optimisation steps are performed:

per-primitive operations :
Since the current (DirectX 11) backend basically offers three stages for geometry processing (vertex shader, tesselation shader, and geometry shader) with different capabilities and associated computation rates our composed modules need to be mapped onto these efficiently. Since the computation rates for these hardware-stages are generally unknown (we don't know what our inputs will look like) we decided to move all operations to the earliest stage possible, with the underlying assumption, that each operation is thus performed at the lowest rate. It is, for example, possible to move geometry shader operations, which are equally performed for all vertices, to the vertex shader. Similar rules can be derived for the other hardware-stages, an overview of these rules can be found in Section 4. A number of additional limitations (due to hardware capabilities) are introduced, e.g, the tessellation-shader needs to be done prior to a geometry-shader, etc.

per-pixel operations :
pixel or fragment-shaders are divided into two parts: the first part represents calculations invariant respective to rasterizer-interpolation, which can therefore be performed per primitive (in general the faster solution). The second part consists of all operations that can only be performed in a pixel shader. This splitting may cause additional traffic for the rasterizer-interpolation since this may need e.g. the interpolation of a vector instead of a scalar. Since these costs are hard to estimate we move only operations where the interpolated type does not exceed the result type in size.

Shader module specialization. Since a single generated shader module might still cover a number of different input-setups (textured vs. non-textured, etc.) using shader-control-flow we provide methods for simply specialising a shader module using contextual information (e.g. there are no textures available, etc.) The shader modules are then partially evaluated using this information and recompiled for the backend. If, for example, a geometry does not contain normals the corresponding shader modules are optimised to elim-

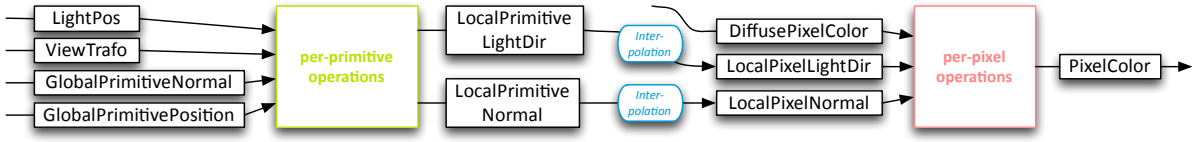

Figure 5: A simple shader module consists of explicitly specified per-primitive and per-pixel operations each with their semantic inputs and outputs. This shows the example from figure 4 tagged with semantic inputs and outputs.

inate any code that accesses normals of the geometry, thereby improving rendering performance.

Shader module unification. Since shaders are compositions of abstract modules it's relatively easy to find common operations for them using the high level information provided by the composition operators. If the rendering performance can be improved by reducing switches between shaders, two shader modules can be unified using simple control flow, adding a parameter to select the shader module as an additional input to the combined shader module.

Common subexpression elimination. Although common subexpression elimination results in optimal code respective to the number of operations, additional temporary variables stressing the *HLSL*-compiler need to be introduced. In optimizing compilers, sophisticated analysis carefully choose subexpressions to be considered for code motion (e.g. (Knoop et al., 1992)) in order to limit temporaries. Our system in contrast heuristically eliminates expressions exceeding a syntactic complexity threshold. These complexities are based on estimated complexities for all intrinsic functions which are simply summed for each expression. With this simple scheme the *HLSL*-compiler does a good job in optimizing shaders while maintaining good compile-time performance.

Constant/Uniform-calculations. All computations resulting in a constant value (for each draw call) can be pre-calculated by the rendering system. Since a brute force approach would result in a large number of uniform-parameters only calculations exceeding a certain complexity (as mentioned above) are considered.

Arithmetic optimizations. Since there are only very few restrictions on how to compose shaders (i.e. outputs and inputs must match), it is possible to introduce unnecessary calculations through these compositions (e.g. `normalize(normalize(vector))`, $(a - a)$, etc.). Similarly to tree parsers (Fraser et al., 1992), used for instruction selection in code generators our optimizer maintains a set of expression patterns with associated rewrite rules and some estimated cost. Notably, our system also considers domain-specific knowledge as a variables vector-basis for further optimization. As an example, `ViewMatrix*ModelMatrix` is transformed to use the uniform `ModelViewMatrix` in order to eliminate expensive matrix multiplications.

User guided simplification. Additional contextual information can be specified for shader inputs values. As an example, the user may annotate the vertex colors to be constant or the normals to be constant per face.

Using these annotations the backend stage mapping can perform further optimisations by moving operations to earlier stages. Together with redundancy removal, dead code elimination, and constant/uniform calculations this can lead to significantly simplified shaders. As an example, it is unnecessary to interpolate face normals in a shader, when the normals are known to be constant per primitive.

Further optimization possibilities. Our abstract pipeline representation is general enough to support completely different approaches like perceptual simplification methods (Sitthi-Amorn et al., 2011), or automatic approaches exploiting temporal coherence (Sitthi-Amorn et al., 2008), which we will pursue in the future.

4 Implementation

Our shading language is implemented as an extension of an existing in-house rendering framework written in C#. Of course it is possible to implement our expression tree based approach with any language that provides abstract data types, the use of anonymous functions significantly reduces the syntactic overhead. A C++ 11 implementation would be equivalent to our approach, while a Java implementation would use anonymous classes instead of anonymous functions.

The expression trees in our approach are created and combined using so called *shader types*, which represent predefined data types available to shaders

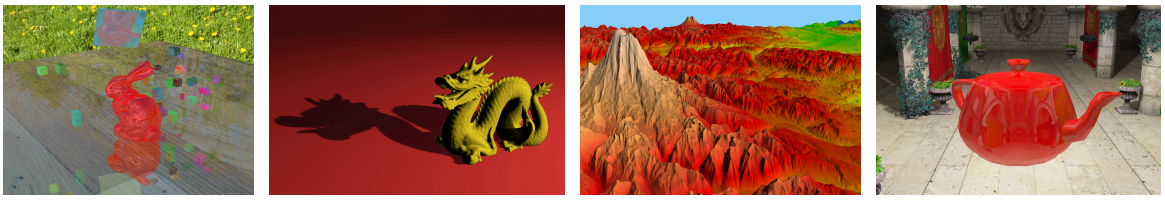

Figure 6: Shader effects composed with our approach. From left to right: raytraced reflections with simple texturing, a composition of an illumination and shadow mapping shader, and a subsivision shader (all from section 5), as well as a screenshot from our lighting design application (see Section 6.1) demonstrating texturing and reflective materials with environment mapping.

(e.g. vectors, matrices, textures, aso.). Each shader type provides methods (e.g. the operators + and – or the dot product) that do not actually perform operations, but build an expression tree for the corresponding operations. Thus each expression that specifies a shader module, returns the complete expression tree for that shader module upon execution.

Due to the flexibility of our functional-style implementation we are not limited to predefined control flow statements such as the RenderMan Shading Language (RSL) (Hanrahan and Lawson, 1990) : we provide higher-order functions that encapsulate conditional evaluation and loops. This makes it possible to integrate conditions and loops into expression trees in typical implementation languages of rendering frameworks, even if it is not possible to overload intrinsic language constructs such as the conditional evaluation operator $condition\ ?\ value_1 : value_2$ and the *for loop* for shader types:

```
var floatVal = Fun.IfThenElse<Float>(c < 1.0f, c, 1.0f);

var initial = new { Index = Int.Zero, Val = Float.Zero };
var diffuse = Fun.Loop(initial, i => i.Index < lightCount,
    i => {
      var dir = light[i.Index] - worldPosition.XYZ;
      return new { Index = i.Index + 1,
                   Val = i.Val + normal.Dot(dir.Normalized) };
    });
```

Figure 7: example usage of higher-order functions for loops and conditional execution

As mentioned in Section 3.4, the evaluation of each expression is moved to the earliest possible stage in any given hardware pipeline. In the following list we give the conditions for performing the indicated optimizations:

PixelShader → GeometryShader : An expression can be moved, if it comprises a linear function. Note that functions can be linear under specific circumstances, e.g. if one function argument is a constant.

GeometryShader → DomainShader/VertexShader : If the same function is applied to all vertex-dependent inputs (i.e. it appears for each of the vertices), the function-expression with its arguments can be moved.

DomainShader → HullShader : If an expression does not contain the tessellation coordinate (i.e. the *domain location*) it can be moved.

DomainShader → HullShaderConstantFunction Similar to *DomainShader → HullShader*.

HullShader → VertexShader : Identical preconditions to GeometryShader → VertexShader stage.

Although we focussed on the Rules for DirectX 11 and OpenGL 4 we also implemented an experimental backend for our OpenCL based raytracer, which only supports one shader-stage computing the color for a primitive at a certain coordinate. Due to our abstract stage interpretation the modules could easily be mapped onto this stage when possible (features like tessellation are currently not supported by the raytracer)

The first step of the compilation process is the creation of a single pipeline for the completely combined shader modules. This abstract pipeline is then processed using the optimisation stages shown in Section 3.4. The output of the optimisation process is a complete hardware shader in the shading language of the backend: in our case a complete HLSL shader for the DirectX 11 backend.

5 Examples

In order to demonstrate the applicability of our approach to common techniques, we provide code excerpts for per-pixel lighting, shadow mapping and subdivision (see Figure 6). Furthermore, we demonstrate the interactive capabilities of our system in the accompanying video.

The following example shows the full implementation of a basic transformation pipeline with per-pixel lighting. Note that user-defined parameters are communicated by name (e.g. `"LightPosition"`), and vertices of primitives supplied to the primitive shader can be accessed using iterators (e.g. `.DoByVertex`). We also provide swizzle operators (e.g. `.XYZ`) including constants (letter `O` is zero, letter `I` is one) for all vector types.

```
public class PerPixelLighting : Module {

  public class Vertex {
    public Float4 Position = Varying.Position;
    public Float3 Normal = Varying.Normal;
    public Float4 WorldPosition = Varying.WorldPosition;
    public Float4 Color = Varying.Color;
  }

  public class Pixel {
    public Float4 Color = Varying.Color;
    public Pixel(Float3 color, Float alpha = 1) {
       Color = new Float4(color, alpha);
    }
  }

  public Fragment<Pixel> Shader(AnyPrimitive<Vertex> input) {
    var transformed = input.DoByVertex(v => {
      v.Position = Uniform.ModelViewProjTrafo * v.Position;
      v.Normal = Uniform.NormalTrafo * v.Normal;
      v.WorldPosition = Uniform.ModelTrafo * v.Position;
      return v;
    });

    return transformed.Rasterize(f => {
      var dir = (Uniform.LightPositions[0]
                 - f.WorldPosition).XYZ.Normalized;
      return new Pixel(dir.Dot(f.Normal.Normalized) * f.Color.XYZ,
                       f.Color.W);
    });
  }
}
```

Figure 8: vertex transformation and per pixel lighting example

Geometry vertex colors are automatically bound to the *Varying.Color* input if needed.

Here we demonstrate the composability of modules by combining per-pixel lighting from Figure 8 with a simple shadow mapping module.

```
public class ShadowMapping : Module {
  public Float4x4 ShadowMapTrafo;
  public Texture2D ShadowMap;

  private Float3 GetShadowTexCoord(Float4 worldPos) {
    var p = ShadowMapTrafo * worldPos; var pp = p.XYZ / p.W;
    var tc = new Float3((Float2.II + pp.XY) * 0.5f, pp.Z);
    tc.Y = 1 - tc.Y; return tc;
  }

  public Fragment<Pixel> Shader(AnyPrimitive<Vertex> input) {
    return input.Rasterize(f => {
      var mytc = GetShadowTexCoord(f.WorldPosition);
      var smValue = ShadowMap.SampleCmp(mytc.XY, mytc.Z);
      return new Pixel(f.Color.XYZ * smValue, f.Color.W);
    });
}}
```

Figure 9: shadow-mapping example

Both modules can be simply composed in the following way:

```
...
var sg = ...                        // some scene graph node
var shadowMapping = new ShadowMapping();
var surface = Composition.Sequence.Compose(new PerPixelLighting(),
                                           shadowMapping);
shadowSurface.ShadowMap = renderTarget.DepthTexture;
shadowSurface.ShadowMapTrafo = ... // the transformation value
sg = sg.Surface(surface);
...
```

Figure 10: composition of lighting and shadow mapping

The code shown in Figure 10 assigns *values* to *module inputs* by implicitly creating uniform inputs in the backend code and setting their values using the renderer infrastructure. These values can thus be changed at runtime.

```
public Float GetFactor(Float4 p0, Float4 p1) {
  var len = (p1.XYZ - p0.XYZ).Length;
  return Float.Clamp(len / MaxLineLength, 1, 64);
}

public Triangle<Vertex> Shader(Triangle<Vertex> input) {
  return input.DoByPrimitive( tri => {
    var f0 = GetFactor(tri.P1.Position, tri.P2.Position);
    var f1 = GetFactor(tri.P2.Position, tri.P0.Position);
    var f2 = GetFactor(tri.P0.Position, tri.P1.Position);
    var factors = new TessellationFactors();
    // tessellation is defined as in OpenGL 4/DirectX 11
    factors.EdgeFactors = new[] { f0, f1, f2 };
    factors.InnerFactors[0] = (f0 + f1 + f2) / 3.0f;
    return factors;
  },
  (tri, constant, crd) => return new Vertex() {
    Position = Float4.Lerp(
               tri.P0.Position, tri.P1.Position, tri.P2.Position,
               crd)
  });
}
```

Figure 11: simple tessellation example

Figure 11 shows a simple subdivision operation mapped to the DirectX 11 tessellation stages. The first anonymous function which calculates the tessellation factors can return any custom type inheriting from `TessellationFactors`. The `constant` argument in the interpolation function then refers to that type. Although it would theoretically be possible to create arbitrary output triangles for a given input patch using the DirectX/OpenGl tessellation stages we decided to expose the functionality as provided by our main backend. The first lambda function basically corresponds to the Hull-/TessellationControl-Shader and the second to the Domain-/TessellationEvaluation-Shader.

6 Analysis

6.1 A real-world comparison

We evaluated our concept by re-implementing all shaders used in a production-quality real-world application for lighting design: it uses shaders for computing global illumination, for drawing lines and points

in debugging and editing views, and for rendering a number of different materials with diffuse, and specular components and environment maps for realistic looking reflections. The complexity of shaders ranges from simple flat-shading all the way to a global-illumination shader that needs to perform polygon clipping for each rendered pixel (see Figure 6 as well as the demonstration in the accompanying video). The following table gives an overview of the number of shader modules and the lines of code for the same functionality using standard HLSL and our composed shader modules (CoSMo).

	lines	# modules	compile time
HLSL	2324	37	7.5s
CoSMo	1148	20	11.8s

Table 1: Code metrics for a real-world lighting design application.

We were able to roughly halve the number of code-lines as well as the number of modules, while maintaining comparable shader-compile-times and shader performance. Due to the high level of reuse enabled by our shader modules, the shader modules implemented for this application are a lot less specific, and can therefore be reused in a number of future applications.

6.2 Performance

We evaluated CoSMo performance using three sample shaders:

skinning: renders two animated simple meshes replicated 50 times without hardware instancing with standard skinning and diffuse lighting, where skinning is performed with a maximum of four bone influences per vertex.

tesselation: subdivides all 871414 triangles of the Stanford dragon to 1-pixel-sized triangles using simple interpolation of positions, normals and light direction, and applies diffuse lighting.

shadow mapping: renders a shadow mapped Standard dragon using a pre-calculated shadow map. A 9x9 Gaussian filter implemented in the pixel shader is used to blur the shadow.

All our tests were performed on an Intel(R) Core(TM) i7 CPU 930 @ 2.80 GHz system with 12 GB of RAM and a GeForce GTX 480 graphics card.

In our first evaluation (see Table 2) we compared the final rendering speed of manually coded shaders with CoSMo-generated shaders. We found that in most cases CoSMo produced equally fast or faster shader code. We rely on the aggressive optimizations in the compiler backend (HLSL or GLSL) to overcome some of the deficiencies in our code generator. The tessellation example was slightly faster since some calculations performed in the DomainShader were automatically moved to the vertex shader.

shader	manual	CoSMo	factor
skinning	35 fps	35 fps	1.00
tesselation	178 fps	189 fps	1.06
shadow-mapping	209 fps	209 fps	1.00

Table 2: Performance results in frames per second (fps) comparing CoSMo generated shaders with manually coded shaders.

Our second evaluation (see Table 3) shows the compile times for manually coded shaders compared to the sum of code generation and compile times for CoSMo shaders. The additional generation of HLSL from CoSMo code results in compile-times that are less than 2 times longer for the given examples, an acceptable overhead given the increased flexibility and reusability.

shader	manual	CoSMo	factor
skinning	140 ms	217 ms (78+139)	1.54
tessellation	46 ms	81 ms (33+48)	1.77
shadow-map	85 ms	110 ms (24+86)	1.30

Table 3: Compile times in milliseconds (ms) of manually written HLSL code and CoSMo code. The CoSMo compile times are given as the sum of generation of HLSL from CoSMo and the HLSL compile time.

7 Conclusion and Future Work

We have demonstrated a powerful framework for combining simple shader modules into complex shaders, by providing semantic composition of modules, an abstract model of two shader stages and a powerful optimizing backend for mapping this programming model onto existing hardware pipelines. With only moderate increase in shader compilation time, this results in a significant reduction of code while retaining the execution performance of hand-coded shaders. Shader composition is performed fast enough for interactively combining shaders, allowing rapid shader development in an explorative manner. Due to the high reusability of the shader modules of our new approach, we have already started to build a comprehensive library of modules that can be freely

combined and provide a framework for rapid development of rendering applications that minimises the necessity of writing shader modules.

However, our current treatment of tessellation modules does not cover all possible composition scenarios and only provides for a single tessellation stage. Thus we cannot handle the sequential composition of tessellation modules. We are currently working on more flexible techniques for specifying and composing tessellation modules, in order to lift this restriction on our semantically specified shaders.

In a number of use cases such as volume rendering and global illumination, we found that using a pure expression based language—although possible—can be somewhat inconvenient, since it requires a complete reformulation of some algorithms that can be easily expressed in procedural languages. We therefore plan to add support for imperative shader fragments while maintaining composability.

ACKNOWLEDGEMENTS

We would like to thank Manuel Wieser for providing 3D models, especially *Eigi, The Dinosaur*. The competence center VRVis is funded by BMVIT, BMWFJ, and City of Vienna (ZIT) within the scope of COMET Competence Centers for Excellent Technologies. The program COMET is managed by FFG.

REFERENCES

Austin, C. A. (2005). Renaissance: A Functional Shading Language. Master's thesis, Iowa State University, Ames, Iowa, USA.

Cook, R. L. (1984). Shade Trees. *SIGGRAPH Comp. Graph.*, 18(3):223–231.

Elliott, C. (2004). Programming Graphics Processors Functionally. In *Proc. of the 2004 ACM SIGPLAN workshop on Haskell*, Haskell '04, pages 45–56, New York, NY, USA. ACM.

Foley, T. and Hanrahan, P. (2011). Spark: Modular, Composable Shaders for Graphics Hardware. *ACM Transactions on Graphics*, 30(4):107:1–107:12.

Fraser, C. W., Henry, R. R., and Proebsting, T. A. (1992). BURG: fast optimal instruction selection and tree parsing. *SIGPLAN Not.*, 27(4):68–76.

Hanrahan, P. and Lawson, J. (1990). A Language for Shading and Lighting Calculations. *SIGGRAPH Comput. Graph.*, 24(4):289–298.

Kessenich, J., Baldwin, D., and Rost, R. (2012). OpenGL Shading Language, version 4.3. Available from: `http://www.opengl.org/documentation/glsl/` (accessed 23 October 2012).

Knoop, J., Rüthing, O., and Steffen, B. (1992). Lazy code motion. *SIGPLAN Not.*, 27(7):224–234.

Kuck, R. and Wesche, G. (2009). A Framework for Object-Oriented Shader Design. In *Proc. of the 5th International Symposium on Advances in Visual Computing: Part I*, ISVC '09, pages 1019–1030, Berlin, Heidelberg. Springer-Verlag.

Mark, W. R., Glanville, R. S., Akeley, K., and Kilgard, M. J. (2003). Cg: a system for programming graphics hardware in a C-like language. *ACM Trans. Graph.*, 22(3):896–907.

McCool, M., Du Toit, S., Popa, T., Chan, B., and Moule, K. (2004). Shader Algebra. In *ACM SIGGRAPH 2004 Papers*, SIGGRAPH '04, pages 787–795, New York, NY, USA. ACM.

McCool, M. D., Qin, Z., and Popa, T. S. (2002). Shader Metaprogramming. In *Proc. of the ACM SIGGRAPH/EUROGRAPHICS conf. on Graph. Hardware*, HWWS '02, pages 57–68, Aire-la-Ville, Switzerland. Eurograph. Assoc.

McCool, M. D. and Toit, S. D. (2004). *Metaprogramming GPUs with Sh.* A K Peters.

McGuire, M. (2005). The SuperShader. Shader X4: Advanced Rendering Techniques, chapter 8.1, pages 485–498. Cengage Learning Emea.

McGuire, M., Stathis, G., Pfister, H., and Krishnamurthi, S. (2006). Abstract Shade Trees. In *Proc. of the 2006 symposium on Interactive 3D graphics and games*, I3D '06, pages 79–86, New York, NY, USA. ACM.

Microsoft (2010). Shader model 5 DirectX HLSL. from: `http://msdn.microsoft.com/en-us/library/windows/desktop/ff471356\%28v=vs.85\%29.aspx` (accessed Oct. 23, 2012).

Microsoft (2012). Programming Guide for HLSL. from: `http://msdn.microsoft.com/en-us/library/bb509635(v=VS.85).aspx` (accessed Oct. 23, 2012).

Perlin, K. (1985). An Image Synthesizer. In *Proc. of the 12th annual conference on Comp. graph. and interactive techniques*, SIGGRAPH '85, pages 287–296, New York, USA. ACM.

Proudfoot, K., Mark, W. R., Tzvetkov, S., and Hanrahan, P. (2001). A Real-time Procedural Shading System for Programmable Graphics Hardware. In *Proc. of the 28th annual conf. on Comp. graph. and interactive techniques*, SIGGRAPH '01, pages 159–170, New York, NY, USA. ACM.

Sitthi-Amorn, P., Lawrence, J., Yang, L., Sander, P. V., Nehab, D., and Xi, J. (2008). Automated Reprojection-Based Pixel Shader Optimization. *ACM Trans. Graph.*, 27(5):127:1–127:11.

Sitthi-Amorn, P., Modly, N., Weimer, W., and Lawrence, J. (2011). Genetic Programming for Shader Simplification. *ACM Trans. Graph.*, 30(6):152:1–152:12.

Trapp, M. and Döllner, J. (2007). Automated Combination of Real-Time Shader Programs. In Cignoni, P. and Sochor, J., editors, *Eurographics 2007 Shortpaper*, pages 53–56. Eurograph. Assoc.